

all_state_capstone_code-Copy1

March 28, 2018

0.0.1 Importing data

To get started two sets of data will be imported and will be analyzed to figure what will be the best way to preprocess them. After running the code below we can see that both data sets have category data that uses letters and continues data that uses numbers. The training data is comprised of 188,318 rows and 132 columns and the test data is comprised of 125,546 rows and 131 columns. The training data has an extra column for loss. This loss is what the insurance company pays out and is what we want to predict. For the purpose of this project we will be using this data set to train and test the model. The algorithm that found will be used to predict the loss value for the test_data set.

```
In [1]: import numpy as np
import pandas as pd
from IPython.display import display
```

```
train_data = pd.read_csv("train.csv")
test_data = pd.read_csv("test.csv")
```

```
display(train_data.head())
print train_data.shape
```

```
display(test_data.head())
print test_data.shape
```

	id	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	...	cont6	\
0	1	A	B	A	B	A	A	A	A	B	...	0.718367	
1	2	A	B	A	A	A	A	A	A	B	...	0.438917	
2	5	A	B	A	A	B	A	A	A	B	...	0.289648	
3	10	B	B	A	B	A	A	A	A	B	...	0.440945	
4	11	A	B	A	B	A	A	A	A	B	...	0.178193	

	cont7	cont8	cont9	cont10	cont11	cont12	cont13	\
0	0.335060	0.30260	0.67135	0.83510	0.569745	0.594646	0.822493	
1	0.436585	0.60087	0.35127	0.43919	0.338312	0.366307	0.611431	
2	0.315545	0.27320	0.26076	0.32446	0.381398	0.373424	0.195709	
3	0.391128	0.31796	0.32128	0.44467	0.327915	0.321570	0.605077	
4	0.247408	0.24564	0.22089	0.21230	0.204687	0.202213	0.246011	

	cont14	loss
0	0.714843	2213.18
1	0.304496	1283.60
2	0.774425	3005.09
3	0.602642	939.85
4	0.432606	2763.85

[5 rows x 132 columns]

(188318, 132)

	id	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	...	cont5	\
0	4	A	B	A	A	A	A	A	A	B	...	0.281143	
1	6	A	B	A	B	A	A	A	A	B	...	0.836443	
2	9	A	B	A	B	B	A	B	A	B	...	0.718531	
3	12	A	A	A	A	B	A	A	A	A	...	0.397069	
4	15	B	A	A	A	A	B	A	A	A	...	0.302678	

	cont6	cont7	cont8	cont9	cont10	cont11	cont12	\
0	0.466591	0.317681	0.61229	0.34365	0.38016	0.377724	0.369858	
1	0.482425	0.443760	0.71330	0.51890	0.60401	0.689039	0.675759	
2	0.212308	0.325779	0.29758	0.34365	0.30529	0.245410	0.241676	
3	0.369930	0.342355	0.40028	0.33237	0.31480	0.348867	0.341872	
4	0.398862	0.391833	0.23688	0.43731	0.50556	0.359572	0.352251	

	cont13	cont14
0	0.704052	0.392562
1	0.453468	0.208045
2	0.258586	0.297232
3	0.592264	0.555955
4	0.301535	0.825823

[5 rows x 131 columns]

(125546, 131)

In [2]: train_data.skew()

Out[2]:

id	-0.002155
cont1	0.516424
cont2	-0.310941
cont3	-0.010002
cont4	0.416096
cont5	0.681622

```
cont6      0.461214
cont7      0.826053
cont8      0.676634
cont9      1.072429
cont10     0.355001
cont11     0.280821
cont12     0.291992
cont13     0.380742
cont14     0.248674
loss       3.794958
dtype: float64
```

0.0.2 Drop data not needed

After looking at the data there are a few changes that need to be made. First the column for "id" needs to be dropped because it does not add any value to our model. Second the "loss" column needs to be separated from the features.

```
In [3]: train_columns = train_data.columns.values
        train_columns = pd.DataFrame(train_columns)

        loss = train_data["loss"]
        ids = train_data["id"]

        train_data_raw = train_data.drop(["id","loss"], axis = 1)

        test_data_id = test_data["id"]
        test_data_raw = test_data.drop(["id"], axis = 1)

        #make sure only ids and loss columns were dropped
        print train_data_raw.shape
        print test_data_raw.shape
```


```
(188318, 130)
```

```
(125546, 130)
```

0.0.3 Loss data

The loss data is plotted on a histogram chart and we can see that the data is skewed to the right so it needs to be normalized. Regression algorithms can be sensitive to the distribution of values and can results in the model underperforming if the data is not normally distributed.

```
In [4]: import matplotlib.pyplot as pl
        pl.hist(loss,bins = 100)
        pl.show()
```




0.0.4 Normalizing the Data

Logarithmic transformation is applied to the "loss" data so that it does not negatively affect the performance of the learning algorithm. Apply the transformation will reduce the range of the values.

```
In [5]: loss_log_transformed = loss.apply(lambda x: np.log(x + 1))

pl.hist(loss_log_transformed,bins = 100)
pl.show()
```



0.0.5 Splitting the data

There are two different types of data in this data set, which are numeric and non-numeric. Below the numeric and non-numeric data will be separated.

```
In [6]: split = 116
        cont_data = train_data_raw.iloc[:,split:]
        display(cont_data.head())
        print cont_data.shape

        split = 116
        cont_test_data = test_data_raw.iloc[:,split:]
        display(cont_test_data.head())
        print cont_test_data.shape
```

	cont1	cont2	cont3	cont4	cont5	cont6	cont7 \
0	0.726300	0.245921	0.187583	0.789639	0.310061	0.718367	0.335060
1	0.330514	0.737068	0.592681	0.614134	0.885834	0.438917	0.436585
2	0.261841	0.358319	0.484196	0.236924	0.397069	0.289648	0.315545
3	0.321594	0.555782	0.527991	0.373816	0.422268	0.440945	0.391128
4	0.273204	0.159990	0.527991	0.473202	0.704268	0.178193	0.247408

	cont8	cont9	cont10	cont11	cont12	cont13	cont14
0	0.30260	0.67135	0.83510	0.569745	0.594646	0.822493	0.714843
1	0.60087	0.35127	0.43919	0.338312	0.366307	0.611431	0.304496
2	0.27320	0.26076	0.32446	0.381398	0.373424	0.195709	0.774425

```

3  0.31796  0.32128  0.44467  0.327915  0.321570  0.605077  0.602642
4  0.24564  0.22089  0.21230  0.204687  0.202213  0.246011  0.432606

```

(188318, 14)

```

      cont1      cont2      cont3      cont4      cont5      cont6      cont7  \
0  0.321594  0.299102  0.246911  0.402922  0.281143  0.466591  0.317681
1  0.634734  0.620805  0.654310  0.946616  0.836443  0.482425  0.443760
2  0.290813  0.737068  0.711159  0.412789  0.718531  0.212308  0.325779
3  0.268622  0.681761  0.592681  0.354893  0.397069  0.369930  0.342355
4  0.553846  0.299102  0.263570  0.696873  0.302678  0.398862  0.391833

      cont8      cont9      cont10      cont11      cont12      cont13      cont14
0  0.61229  0.34365  0.38016  0.377724  0.369858  0.704052  0.392562
1  0.71330  0.51890  0.60401  0.689039  0.675759  0.453468  0.208045
2  0.29758  0.34365  0.30529  0.245410  0.241676  0.258586  0.297232
3  0.40028  0.33237  0.31480  0.348867  0.341872  0.592264  0.555955
4  0.23688  0.43731  0.50556  0.359572  0.352251  0.301535  0.825823

```

(125546, 14)

0.0.6 Numeric Data

Histograms are plotted to what is going on within the continuous data. None of the categories seem to have a normal distribution

In [7]: *#plot continues data*















```

cont_data_columns = cont_data.columns.values
fig, axarr = pl.subplots(nrows=7, ncols=2,figsize=(20, 40))

#columns
index_arr = -1
for i in range(7):
    for j in range(2):
        index_arr = index_arr + 1
        axarr[i, j].hist(cont_data[cont_data_columns[index_arr]], bins = 300)
        axarr[i, j].set_title(cont_data_columns[index_arr])

pl.show()

```



0.0.7 Hot key encoding

Typically learning algorithms expect input to be numeric so all non-numeric data must be converted into numeric data. There are many ways to do so but for this project we will be using one-hot encoding. This method creates dummy variables for each possible category of each non-numeric feature.

```
In [8]: cat_data = train_data_raw.iloc[:, :split]
        display(cat_data.head())
        print cat_data.shape

        cat_test_data = test_data_raw.iloc[:, :split]
        display(cat_test_data.head())
        print cat_test_data.shape
```

	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	cat10	...	cat107	cat108	\
0	A	B	A	B	A	A	A	A	B	A	...	J	G	
1	A	B	A	A	A	A	A	A	B	B	...	K	K	
2	A	B	A	A	B	A	A	A	B	B	...	F	A	
3	B	B	A	B	A	A	A	A	B	A	...	K	K	
4	A	B	A	B	A	A	A	A	B	B	...	G	B	

	cat109	cat110	cat111	cat112	cat113	cat114	cat115	cat116
0	BU	BC	C	AS	S	A	O	LB
1	BI	CQ	A	AV	BM	A	O	DP
2	AB	DK	A	C	AF	A	I	GK
3	BI	CS	C	N	AE	A	O	DJ
4	H	C	C	Y	BM	A	K	CK

[5 rows x 116 columns]

(188318, 116)

	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	cat10	...	cat107	cat108	\
0	A	B	A	A	A	A	A	A	B	A	...	L	K	
1	A	B	A	B	A	A	A	A	B	A	...	F	B	
2	A	B	A	B	B	A	B	A	B	B	...	G	A	
3	A	A	A	A	B	A	A	A	A	A	...	K	K	
4	B	A	A	A	A	B	A	A	A	A	...	E	B	

	cat109	cat110	cat111	cat112	cat113	cat114	cat115	cat116
0	BI	BC	A	J	AX	A	Q	HG
1	BI	CO	E	G	X	A	L	HK
2	BI	CS	C	U	AE	A	K	CK

3	BI	CR	A	AY	AJ	A	P	DJ
4	AB	EG	A	E	I	C	J	HA

[5 rows x 116 columns]

(125546, 116)

```
In [9]: len(np.unique(cat_data['cat1']))
        unique_var = []
        for column in cat_data.columns.values:
            variables = len(np.unique(cat_data[column]))
            unique_var.append(variables)

        # print "{} unique lables in {}".format(variables,column)

In [10]: unique_var2 = list(set(unique_var))
         sorted_var = np.sort(unique_var2)
         for uniquer in sorted_var:
             n = unique_var.count(uniquer)
             print "There is {} unique categories in {} of the lables".format(uniquer,n)
```

```
There is 2 unique categories in 72 of the lables
There is 3 unique categories in 4 of the lables
There is 4 unique categories in 12 of the lables
There is 5 unique categories in 3 of the lables
There is 7 unique categories in 4 of the lables
There is 8 unique categories in 3 of the lables
There is 9 unique categories in 1 of the lables
There is 11 unique categories in 1 of the lables
There is 13 unique categories in 1 of the lables
There is 15 unique categories in 1 of the lables
There is 16 unique categories in 2 of the lables
There is 17 unique categories in 2 of the lables
There is 19 unique categories in 2 of the lables
There is 20 unique categories in 2 of the lables
There is 23 unique categories in 1 of the lables
There is 51 unique categories in 1 of the lables
There is 61 unique categories in 1 of the lables
There is 84 unique categories in 1 of the lables
There is 131 unique categories in 1 of the lables
There is 326 unique categories in 1 of the lables
```

```
In [11]: cat_data_df = cat_data.apply(lambda x: pd.factorize(x)[0])

        cat_test_data_df = cat_test_data.apply(lambda x: pd.factorize(x)[0])

        display(cat_data_df.head())
```

	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	cat10	...	\
0	0	0	0	0	0	0	0	0	0	0	...	
1	0	0	0	1	0	0	0	0	0	1	...	
2	0	0	0	1	1	0	0	0	0	1	...	
3	1	0	0	0	0	0	0	0	0	0	...	
4	0	0	0	0	0	0	0	0	0	1	...	

	cat107	cat108	cat109	cat110	cat111	cat112	cat113	cat114	cat115	\
0	0	0	0	0	0	0	0	0	0	
1	1	1	1	1	1	1	1	0	0	
2	2	2	2	2	1	2	2	0	1	
3	1	1	1	3	0	3	3	0	0	
4	3	3	3	4	0	4	1	0	2	

	cat116
0	0
1	1
2	2
3	3
4	4

[5 rows x 116 columns]

0.0.8 Shuffle and split data

Now that the data has been processed the training data will be split into testing and training in which 80% is used for training and 20% is used for testing.

```
In [12]: cat_cont_data = pd.concat([cat_data_df, cont_data], axis=1)
         display(cat_cont_data.head())
```

```
cat_cont_test_data = pd.concat([cat_test_data_df, cont_test_data], axis=1)
display(cat_cont_test_data.head())
```

	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	cat10	...	\
0	0	0	0	0	0	0	0	0	0	0	...	
1	0	0	0	1	0	0	0	0	0	1	...	
2	0	0	0	1	1	0	0	0	0	1	...	
3	1	0	0	0	0	0	0	0	0	0	...	
4	0	0	0	0	0	0	0	0	0	1	...	

	cont5	cont6	cont7	cont8	cont9	cont10	cont11	\
0	0.310061	0.718367	0.335060	0.30260	0.67135	0.83510	0.569745	
1	0.885834	0.438917	0.436585	0.60087	0.35127	0.43919	0.338312	
2	0.397069	0.289648	0.315545	0.27320	0.26076	0.32446	0.381398	
3	0.422268	0.440945	0.391128	0.31796	0.32128	0.44467	0.327915	

```
4 0.704268 0.178193 0.247408 0.24564 0.22089 0.21230 0.204687
```

```

      cont12  cont13  cont14
0 0.594646 0.822493 0.714843
1 0.366307 0.611431 0.304496
2 0.373424 0.195709 0.774425
3 0.321570 0.605077 0.602642
4 0.202213 0.246011 0.432606
```

[5 rows x 130 columns]

```

      cat1  cat2  cat3  cat4  cat5  cat6  cat7  cat8  cat9  cat10  ...  \
0      0      0      0      0      0      0      0      0      0      0      ...
1      0      0      0      1      0      0      0      0      0      0      ...
2      0      0      0      1      1      0      1      0      0      1      ...
3      0      1      0      0      1      0      0      0      1      0      ...
4      1      1      0      0      0      1      0      0      1      0      ...
```

```

      cont5  cont6  cont7  cont8  cont9  cont10  cont11  \
0 0.281143 0.466591 0.317681 0.61229 0.34365 0.38016 0.377724
1 0.836443 0.482425 0.443760 0.71330 0.51890 0.60401 0.689039
2 0.718531 0.212308 0.325779 0.29758 0.34365 0.30529 0.245410
3 0.397069 0.369930 0.342355 0.40028 0.33237 0.31480 0.348867
4 0.302678 0.398862 0.391833 0.23688 0.43731 0.50556 0.359572
```

```

      cont12  cont13  cont14
0 0.369858 0.704052 0.392562
1 0.675759 0.453468 0.208045
2 0.241676 0.258586 0.297232
3 0.341872 0.592264 0.555955
4 0.352251 0.301535 0.825823
```

[5 rows x 130 columns]

```
In [13]: from sklearn.cross_validation import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(cat_cont_data, loss_log_transformed
```

```
print "Training set has {} samples.".format(X_train.shape[0])
```

```
print "Testing set has {} samples.".format(X_test.shape[0])
```

```
/Users/josegarcia/anaconda2/lib/python2.7/site-packages/sklearn/cross_validation.py:41: Deprecat
"This module will be removed in 0.20.", DeprecationWarning)
```

```
Training set has 141238 samples.
```

```
Testing set has 47080 samples.
```

The code block below will allow for automatically testing different repressors using different data sizes.

```
In [14]: from sklearn.metrics import mean_absolute_error
         from time import time

         def train_predict(learner, sample_size, X_train, y_train, X_test, y_test):
             results = {}
             start = time()
             learner = learner.fit(X_train[:sample_size], y_train[:sample_size])
             end = time()

             results['train_time'] = end - start

             start = time()
             predictions_test = learner.predict(X_test)
             predictions_train = learner.predict(X_train[:300])
             end = time()

             results['pred_time'] = end - start

             results['acc_train'] = mean_absolute_error(np.exp(y_train[:300]), np.exp(predictions_train))

             results['acc_test'] = mean_absolute_error(np.exp(y_test), np.exp(predictions_test))

             print "{} trained on {} samples.".format(learner.__class__.__name__, sample_size)
             print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'], results['acc_test'])
             print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'], results['pred_time'])

             return results
```

0.0.9 Learning Algorithms

Three different repressors will be taken into consideration linear regression, random forest regression, and extreme gradient regression. The mean absolute error of a un-tuned Linear regression model will be used as our bench mark that we need to beat. Mean absolute error is a measure of the average difference between the actual loss and predicted loss.

```
In [15]: from sklearn.linear_model import LinearRegression
         from sklearn.ensemble import RandomForestRegressor
         from xgboost import XGBRegressor

         clf_A = LinearRegression()
         #clf_B = RandomForestRegressor(random_state = 10)
         clf_C = XGBRegressor()
```

```

sample_1 = len(X_train) / 100
sample_10 = len(X_train) / 10
sample_100 = len(X_train) / 1

results = {}

for clf in [clf_A, clf_C]:
    clf_name = clf.__class__.__name__
    results[clf_name] = {}
    for i, samples in enumerate([sample_1, sample_10, sample_100]):
        results[clf_name][i] = train_predict(clf, samples, X_train, y_train, X_test, y_

```

```

LinearRegression trained on 1412 samples.
Trained score = 1349.70 and test score = 1401.50
Trained time = 0.03 and pred time = 0.05
LinearRegression trained on 14123 samples.
Trained score = 1400.24 and test score = 1338.21
Trained time = 0.11 and pred time = 0.03
LinearRegression trained on 141238 samples.
Trained score = 1423.36 and test score = 1320.10
Trained time = 1.36 and pred time = 0.02
XGBRegressor trained on 1412 samples.
Trained score = 1140.34 and test score = 1283.66
Trained time = 0.72 and pred time = 0.27
XGBRegressor trained on 14123 samples.
Trained score = 1320.30 and test score = 1213.91
Trained time = 7.01 and pred time = 0.32
XGBRegressor trained on 141238 samples.
Trained score = 1345.58 and test score = 1205.96
Trained time = 76.89 and pred time = 0.22

```

```

In [16]: results = {}
         start = time()
         learner = XGBRegressor().fit(X_train, y_train)
         end = time()

         results['train_time'] = end - start

         start = time()
         predictions_test = learner.predict(X_test)
         predictions_train = learner.predict(X_train)
         end = time()

         results['pred_time'] = end - start

```

```
results['acc_train'] = mean_absolute_error(np.exp(y_train), np.exp(predictions_train))


results['acc_test'] = mean_absolute_error(np.exp(y_test), np.exp(predictions_test))

print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'], results['acc_test'])
print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'], results['pred_time'])
```

Trained score = 1218.83 and test score = 1205.96
Trained time = 83.07 and pred time = 1.23

```
In [40]: ax = xgb.plot_importance(learner)

fig = ax.figure
fig.set_size_inches(10, 15)
pl.show()
```



```
In [18]: import xgboost as xgb
```

```

dtrain = xgb.DMatrix(data=X_train, label=y_train)
dtest = xgb.DMatrix(data=X_test, label=y_test)

In [19]: from sklearn.metrics import mean_absolute_error
         from time import time
         from sklearn.metrics import make_scorer

xgb_params = {
    'objective': 'reg:linear',
}

scorer = make_scorer(mean_absolute_error)

results = {}
start = time()
bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain, num_boost_round=50,
                    feval=scorer, maximize=False)
end = time()
results['train_time'] = end - start

start = time()
dpred_train = xgb.DMatrix(X_train)
dpred_test = xgb.DMatrix(X_test)

predictions_train = bst_cv1.predict(dtrain)
predictions_test = bst_cv1.predict(dtest)
end = time()
results['pred_time'] = end - start

y_tr = dtrain.get_label()
y_te = dtest.get_label()

results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
print "params {}".format(xgb_params)
print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'], results['acc_test'])
print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'], results['pred_time'])

params {'objective': 'reg:linear'}
Trained score = 1117.33 and test score = 1157.39
Trained time = 74.03 and pred time = 0.84

In [21]: from sklearn.metrics import mean_absolute_error

```



```

from time import time
from sklearn.metrics import make_scorer

xgb_params = {'objective': 'reg:linear'}

scorer = make_scorer(mean_absolute_error)
eta_trials = [ 0.2, 0.3, 0.4, 0.5, 0.6]

for etas in eta_trials:
    xgb_params['eta'] = etas
    results = {}
    start = time()
    bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain, num_boost_round=50,
                        feval=scorer, maximize=False)
    end = time()
    results['train_time'] = end - start

    start = time()
    dpred_train = xgb.DMatrix(X_train)
    dpred_test = xgb.DMatrix(X_test)

    predictions_train = bst_cv1.predict(dtrain)
    predictions_test = bst_cv1.predict(dtest)
    end = time()
    results['pred_time'] = end - start

    y_tr = dtrain.get_label()
    y_te = dtest.get_label()

    results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

    results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
    print "params {}".format(xgb_params)
    print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'],
                                                                    results['acc_test'])
    print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'],
                                                                    results['pred_time'])

params {'objective': 'reg:linear', 'eta': 0.2}
Trained score = 1132.61 and test score = 1164.73
Trained time = 84.28 and pred time = 0.97
params {'objective': 'reg:linear', 'eta': 0.3}
Trained score = 1112.43 and test score = 1162.37
Trained time = 76.11 and pred time = 0.98
params {'objective': 'reg:linear', 'eta': 0.4}
Trained score = 1103.95 and test score = 1174.69
Trained time = 76.48 and pred time = 0.86
params {'objective': 'reg:linear', 'eta': 0.5}

```

Trained score = 1097.96 and test score = 1181.80
Trained time = 73.79 and pred time = 0.87
params {'objective': 'reg:linear', 'eta': 0.6}
Trained score = 1093.13 and test score = 1191.07
Trained time = 84.34 and pred time = 1.22

```
In [22]: from sklearn.metrics import mean_absolute_error
         from time import time
         from sklearn.metrics import make_scorer

         xgb_params = { 'eta': 0.3,
                        'objective': 'reg:linear'}

         scorer = make_scorer(mean_absolute_error)

         max_depth = [6,8,10,12,14]
         child_weight = [4,6,8,10,12]

         for depth in max_depth:
             xgb_params['max_depth'] = depth
             for weight in child_weight:
                 xgb_params['min_child_weight'] = weight
                 results = {}
                 start = time()
                 bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain, num_boost_round=50,
                                     feval=scorer, maximize=False)
                 end = time()
                 results['train_time'] = end - start

                 start = time()

                 predictions_train = bst_cv1.predict(dtrain)
                 predictions_test = bst_cv1.predict(dtest)
                 end = time()
                 results['pred_time'] = end - start

                 y_tr = dtrain.get_label()
                 y_te = dtest.get_label()

                 results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

                 results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
                 print "params {}".format(xgb_params)
                 print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'], results['acc_test'])
                 print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'], results['pred_time'])
```

```

params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 4}
Trained score = 1112.88 and test score = 1164.97
Trained time = 74.89 and pred time = 0.22
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 6}
Trained score = 1112.19 and test score = 1162.07
Trained time = 71.08 and pred time = 0.15
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1114.88 and test score = 1160.58
Trained time = 74.90 and pred time = 0.16
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 10}
Trained score = 1113.64 and test score = 1161.35
Trained time = 69.79 and pred time = 0.15
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 12}
Trained score = 1116.00 and test score = 1163.05
Trained time = 79.54 and pred time = 0.16
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 8, 'min_child_weight': 4}
Trained score = 1033.52 and test score = 1166.14
Trained time = 110.02 and pred time = 0.22
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 8, 'min_child_weight': 6}
Trained score = 1029.29 and test score = 1168.36
Trained time = 108.27 and pred time = 0.23
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 8, 'min_child_weight': 8}
Trained score = 1036.61 and test score = 1165.43
Trained time = 100.98 and pred time = 0.21
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 8, 'min_child_weight': 10}
Trained score = 1036.64 and test score = 1170.84
Trained time = 105.34 and pred time = 0.21
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 8, 'min_child_weight': 12}
Trained score = 1046.63 and test score = 1165.15
Trained time = 101.37 and pred time = 0.22
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 10, 'min_child_weight': 4}
Trained score = 900.37 and test score = 1186.86
Trained time = 139.10 and pred time = 0.41
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 10, 'min_child_weight': 6}
Trained score = 918.06 and test score = 1175.81
Trained time = 135.95 and pred time = 0.28
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 10, 'min_child_weight': 8}
Trained score = 928.50 and test score = 1182.07
Trained time = 123.93 and pred time = 0.30
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 10, 'min_child_weight': 10}
Trained score = 942.46 and test score = 1180.68
Trained time = 122.31 and pred time = 0.29
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 10, 'min_child_weight': 12}
Trained score = 940.81 and test score = 1177.83
Trained time = 122.35 and pred time = 0.34
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 12, 'min_child_weight': 4}
Trained score = 740.77 and test score = 1204.35
Trained time = 160.24 and pred time = 0.55

```

```

params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 12, 'min_child_weight': 6}
Trained score = 769.37 and test score = 1202.26
Trained time = 161.36 and pred time = 0.37
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 12, 'min_child_weight': 8}
Trained score = 802.14 and test score = 1199.07
Trained time = 161.10 and pred time = 0.36
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 12, 'min_child_weight': 10}
Trained score = 805.89 and test score = 1197.55
Trained time = 136.96 and pred time = 0.36
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 12, 'min_child_weight': 12}
Trained score = 834.25 and test score = 1192.50
Trained time = 135.64 and pred time = 0.35
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 14, 'min_child_weight': 4}
Trained score = 550.42 and test score = 1232.94
Trained time = 164.68 and pred time = 0.49
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 14, 'min_child_weight': 6}
Trained score = 603.65 and test score = 1225.92
Trained time = 160.25 and pred time = 0.48
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 14, 'min_child_weight': 8}
Trained score = 632.69 and test score = 1218.75
Trained time = 161.40 and pred time = 0.45
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 14, 'min_child_weight': 10}
Trained score = 683.73 and test score = 1216.06
Trained time = 157.83 and pred time = 0.43
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 14, 'min_child_weight': 12}
Trained score = 702.02 and test score = 1214.76
Trained time = 157.19 and pred time = 0.43

```

```

In [23]: from sklearn.metrics import mean_absolute_error
         from time import time
         from sklearn.metrics import make_scorer

```

```

xgb_params = {'eta': 0.3,
              'objective': 'reg:linear',
              'max_depth': 6,
              'min_child_weight': 8,
              }

```

```

scorer = make_scorer(mean_absolute_error)
eta_trials = [0.05, 0.1, 0.2, 0.3]

```

```

for etas in eta_trials:
    xgb_params['eta'] = etas
    results = {}
    start = time()

```

```

bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain, num_boost_round=100,
                    feval=scorer, maximize=False)

end = time()
results['train_time'] = end - start

start = time()
dpred_train = xgb.DMatrix(X_train)
dpred_test = xgb.DMatrix(X_test)

predictions_train = bst_cv1.predict(dtrain)
predictions_test = bst_cv1.predict(dtest)
end = time()
results['pred_time'] = end - start

y_tr = dtrain.get_label()
y_te = dtest.get_label()

results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
print "params {}".format(xgb_params)
print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'],
results['acc_test'])
print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'],
results['pred_time'])

```

```

params {'objective': 'reg:linear', 'eta': 0.05, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1193.83 and test score = 1209.61
Trained time = 135.15 and pred time = 1.16
params {'objective': 'reg:linear', 'eta': 0.1, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1127.89 and test score = 1158.75
Trained time = 131.56 and pred time = 0.95
params {'objective': 'reg:linear', 'eta': 0.2, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1093.88 and test score = 1152.22
Trained time = 136.25 and pred time = 0.95
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1074.92 and test score = 1157.44
Trained time = 137.87 and pred time = 0.87

```

```

In [24]: from sklearn.metrics import mean_absolute_error
        from time import time
        from sklearn.metrics import make_scorer

```

```

xgb_params = {'eta': 0.1,
              'objective': 'reg:linear',
              'max_depth': 6,
              'min_child_weight': 8,

```

```

}

scorer = make_scorer(mean_absolute_error)
eta_trials = [0.05, 0.1 , 0.2, 0.3]

for etas in eta_trials:
    xgb_params['eta'] = etas
    results = {}
    start = time()
    bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain, num_boost_round=200,
                        feval=scorer, maximize=False)
    end = time()
    results['train_time'] = end - start

    start = time()
    dpred_train = xgb.DMatrix(X_train)
    dpred_test = xgb.DMatrix(X_test)

    predictions_train = bst_cv1.predict(dtrain)
    predictions_test = bst_cv1.predict(dtest)
    end = time()
    results['pred_time'] = end - start

    y_tr = dtrain.get_label()
    y_te = dtest.get_label()

    results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

    results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
    print "params {}".format(xgb_params)
    print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'],
    print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'],r

params {'objective': 'reg:linear', 'eta': 0.05, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1127.56 and test score = 1160.22
Trained time = 262.71 and pred time = 1.33
params {'objective': 'reg:linear', 'eta': 0.1, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1089.44 and test score = 1147.20
Trained time = 265.51 and pred time = 1.30
params {'objective': 'reg:linear', 'eta': 0.2, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1046.76 and test score = 1151.11
Trained time = 268.97 and pred time = 1.32
params {'objective': 'reg:linear', 'eta': 0.3, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1019.84 and test score = 1165.08
Trained time = 262.76 and pred time = 1.25

```

```

In [49]: X_train2, X_test2, y_train2, y_test2 = train_test_split(cat_cont_data, loss_log_transfo

dtrain2 = xgb.DMatrix(data=X_train2, label=y_train2)
dtest2 = xgb.DMatrix(data=X_test2, label=y_test2)

In [50]: from sklearn.metrics import mean_absolute_error
from time import time
from sklearn.metrics import make_scorer

xgb_params = {'eta': 0.1,
              'objective': 'reg:linear',
              'max_depth': 6,
              'min_child_weight': 8,
              }

scorer = make_scorer(mean_absolute_error)

results = {}
start = time()
bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain2, num_boost_round=200,
                    feval=scorer, maximize=False)

end = time()
results['train_time'] = end - start

start = time()

predictions_train = bst_cv1.predict(dtrain2)
predictions_test = bst_cv1.predict(dtest2)
end = time()
results['pred_time'] = end - start

y_tr = dtrain2.get_label()
y_te = dtest2.get_label()

results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
print "params {}".format(xgb_params)
print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'], results['acc_test'])
print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'], results['pred_time'])

params {'objective': 'reg:linear', 'eta': 0.1, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1088.67 and test score = 1146.85
Trained time = 273.97 and pred time = 0.68

In [53]: X_train2, X_test2, y_train2, y_test2 = train_test_split(cat_cont_data, loss_log_transfo

```

```

dtrain2 = xgb.DMatrix(data=X_train2, label=y_train2)
dtest2 = xgb.DMatrix(data=X_test2, label=y_test2)

In [54]: from sklearn.metrics import mean_absolute_error
         from time import time
         from sklearn.metrics import make_scorer

xgb_params = {'eta': 0.1,
              'objective': 'reg:linear',
              'max_depth': 6,
              'min_child_weight': 8,
              }

scorer = make_scorer(mean_absolute_error)

results = {}
start = time()
bst_cv1 = xgb.train(params=xgb_params, dtrain=dtrain2, num_boost_round=200,
                    feval=scorer, maximize=False)

end = time()
results['train_time'] = end - start

start = time()

predictions_train = bst_cv1.predict(dtrain2)
predictions_test = bst_cv1.predict(dtest2)
end = time()
results['pred_time'] = end - start

y_tr = dtrain2.get_label()
y_te = dtest2.get_label()

results['acc_train'] = mean_absolute_error(np.exp(y_tr), np.exp(predictions_train))

results['acc_test'] = mean_absolute_error(np.exp(y_te), np.exp(predictions_test))
print "params {}".format(xgb_params)
print "Trained score = {:.2f} and test score = {:.2f}".format(results['acc_train'], results['acc_test'])
print "Trained time = {:.2f} and pred time = {:.2f}".format(results['train_time'], results['pred_time'])

params {'objective': 'reg:linear', 'eta': 0.1, 'max_depth': 6, 'min_child_weight': 8}
Trained score = 1085.92 and test score = 1151.03
Trained time = 278.26 and pred time = 0.82

```


0.0.10 Conclusion

After attempting to tune the model using GridSearchCV the un-tuned model mean absolute error 1261 and the tuned model had a mean absolute error of 1312. The goal is to have the lowest mean absolute error so we will have to go with the un-tuned model. There is still many ways of improving the model that can be attempted in the future. To start off with because there are so many features in the data one can improve the accuracy of the model by finding which features are most important. Along with this correlation between features is a good way to visualize how the different features are interacting with each other.

```
In [25]: y_dummy = range(1,125547)
```

```
    d = {'loss': y_dummy}
    df = pd.DataFrame(data=d)
```

```
In [26]: dtest_final = xgb.DMatrix(data=cat_cont_test_data,label=y_dummy)
```

```
In [27]: final_predictions = bst_cv1.predict(dtest_final)
        test_data_id
        print final_predictions.shape
        print test_data_id.shape
```

```
(125546,)
```

```
(125546,)
```

```
In [28]: d = {'id': test_data_id, 'loss': final_predictions}
        df = pd.DataFrame(data=d)
        loss_final = np.exp(df['loss'])
        print
        df_final = pd.concat([test_data_id, loss_final], axis=1)
```

```
display(df_final.head())
```

	id	loss
0	4	1397.102905
1	6	1865.396729
2	9	5971.495605
3	12	4782.812500
4	15	1503.430176

```
In [29]: df_final.to_csv('submission.csv')
```